

Dataset Run Report

Current Parsed Source

Source file parsed:

`/Users/pod/Desktop/goodreads_reviews_spoiler.json`

This file is the UCSD Goodreads spoiler-review JSONL file. It stores each record as one JSON object per line with review text split across `review_sentences`. The parser now flattens those sentence lists into one full-review text per record.

Important caveat: this file is **not the dedicated Young Adult genre subset** by itself. If the final project must explicitly be Young Adult, use `goodreads_reviews_young_adult.json.gz` or join these review `book_id` values to Goodreads book/genre metadata and filter to YA. The current files are still valid Goodreads review-text artifacts for the pipeline, but the write-up should not call them Young Adult unless that genre filtering is completed.

Parse Summary

Parsed on: 2026-04-26

Sampling seed: 20260426

Raw records scanned: 1,378,033

Usable records before sampling: 1,373,010

Sample requested: 16,200

Sample written: 16,200

The sample was stratified by rating and review length bucket. The extra 200 rows were collected so that exact target sizes remained possible after duplicate removal.

Cleaning and Split Summary

Rows after quality filter: 16,200

Duplicate reviews removed: 83

Rows available after dedupe: 16,117

Final locked files:

- `data/processed/train.csv`: 12,000 rows
- `data/processed/test.csv`: 3,000 rows
- `data/processed/holdout_locked.csv`: 1,000 rows

This satisfies the assignment target of 15,000 train/test texts plus a 1,000-text locked holdout.

Rating Distribution After Cleaning

Rating	Count
0 / missing	2,685
1	2,695
2	2,687
3	2,683
4	2,676
5	2,691

Length Distribution After Cleaning

Length Bucket	Count
short	5,321
medium	5,398
long	5,398

Word Count Summary

Statistic	Value
Mean	162.2
Std. dev.	221.6
Min	3
25th percentile	22
Median	74
75th percentile	214
Max	2,503

Manual Work Minimization

Created artifacts:

- `data/labels/holdout_labeling_workbook.xlsx`: Excel workbook with instructions, pilot sheet, holdout sheet, coder columns, dropdowns, frozen panes, and wrapped review text.
- `data/labels/holdout_human_labeling_sheet.csv`: script-friendly labeling CSV.
- `data/results/dictionary_baseline_predictions.csv`: automatic dictionary baseline for comparison, not a replacement for human holdout labels.

Recommended human workflow:

1. Three coders label only the 60-row Pilot sheet first.
2. Meet once to resolve codebook confusion.
3. Then all three coders independently label the Holdout sheet.
4. Export/save the completed workbook as CSV if needed, then run `scripts/reliability_report.py`.